

Linux Cross Reference

Free Electrons

Embedded Linux Experts

• [Source Navigation](#) • [Diff Markup](#) • [Identifier Search](#) • [Freetext Search](#) •

Version: 2.0.40 2.2.26 2.4.37 3.11 3.12 3.13 3.14 3.15 3.16 3.17 3.18 3.19 4.0 4.1 4.2 4.3 4.4 4.5 4.6 4.7 4.8

Linux/net/sched/sch_red.c

```

1  /*
2  * net/sched/sch_red.c  Random Early Detection queue.
3  *
4  *      This program is free software; you can redistribute it and/or
5  *      modify it under the terms of the GNU General Public License
6  *      as published by the Free Software Foundation; either version
7  *      2 of the License, or (at your option) any later version.
8  *
9  * Authors:   Alexey Kuznetsov, <kuznet@ms2.inr.ac.ru>
10 *
11 * Changes:
12 * J Hadi Salim 980914: computation fixes
13 * Alexey Makarenko <makar@phoenix.kharkov.ua> 990814: gave on idle link was calculated incorrectly.
14 * J Hadi Salim 980816: ECN support
15 */
16
17 #include <linux/module.h>
18 #include <linux/types.h>
19 #include <linux/kernel.h>
20 #include <linux/skbuff.h>
21 #include <net/pkt_sched.h>
22 #include <net/inet_ecn.h>
23 #include <net/red.h>
24
25
26 /*      Parameters, settable by user:
27 -----
28
29      limit          - bytes (must be > qth_max + burst)
30
31      Hard limit on queue length, should be chosen >qth_max
32      to allow packet bursts. This parameter does not
33      affect the algorithms behaviour and can be chosen
34      arbitrarily high (well, less than ram size)
35      Really, this limit will never be reached
36      if RED works correctly.
37 */
38
39 struct red_sched_data {
40     u32          limit;          /* HARD maximal queue length */
41     unsigned char flags;
42     struct timer_list adapt_timer;
43     struct red_parms parms;
44     struct red_vars vars;
45     struct red_stats stats;
46     struct Qdisc *qdisc;
47 };
48
49 static inline int red_use_ecn(struct red_sched_data *q)
50 {
51     return q->flags & TC_RED_ECN;
52 }
53
54 static inline int red_use_harddrop(struct red_sched_data *q)
55 {
56     return q->flags & TC_RED_HARDDROP;
57 }
58
59 static int red_enqueue(struct sk_buff *skb, struct Qdisc *sch,
60                       struct sk_buff **to_free)
61 {
62     struct red_sched_data *q = qdisc_priv(sch);
63     struct Qdisc *child = q->qdisc;
64     int ret;
65
66     q->vars.qavg = red_calc_qavg(&q->parms,
67                                &q->vars,
68                                child->qstats.backlog);
69
70     if (red_is_idling(&q->vars))
71         red_end_of_idle_period(&q->vars);
72
73     switch (red_action(&q->parms, &q->vars, q->vars.qavg)) {
74     case RED_DONT_MARK:
75         break;
76

```

```

77     case RED_PROB_MARK:
78         qdisc_qstats_overlimit(sch);
79         if (!red_use_ecn(q) || !INET_ECN_set_ce(skb)) {
80             q->stats.prob_drop++;
81             goto congestion_drop;
82         }
83
84         q->stats.prob_mark++;
85         break;
86
87     case RED_HARD_MARK:
88         qdisc_qstats_overlimit(sch);
89         if (red_use_harddrop(q) || !red_use_ecn(q) ||
90             !INET_ECN_set_ce(skb)) {
91             q->stats.forced_drop++;
92             goto congestion_drop;
93         }
94
95         q->stats.forced_mark++;
96         break;
97     }
98
99     ret = qdisc_enqueue(skb, child, to_free);
100     if (likely(ret == NET_XMIT_SUCCESS)) {
101         qdisc_qstats_backlog_inc(sch, skb);
102         sch->q.qlen++;
103     } else if (net_xmit_drop_count(ret)) {
104         q->stats.pdrop++;
105         qdisc_qstats_drop(sch);
106     }
107     return ret;
108
109 congestion_drop:
110     qdisc_drop(skb, sch, to_free);
111     return NET_XMIT_CN;
112 }
113
114 static struct sk_buff *red_dequeue(struct Qdisc *sch)
115 {
116     struct sk_buff *skb;
117     struct red_sched_data *q = qdisc_priv(sch);
118     struct Qdisc *child = q->qdisc;
119
120     skb = child->dequeue(child);
121     if (skb) {
122         qdisc_bstats_update(sch, skb);
123         qdisc_qstats_backlog_dec(sch, skb);
124         sch->q.qlen--;
125     } else {
126         if (!red_is_idling(&q->vars))
127             red_start_of_idle_period(&q->vars);
128     }
129     return skb;
130 }
131
132 static struct sk_buff *red_peek(struct Qdisc *sch)
133 {
134     struct red_sched_data *q = qdisc_priv(sch);
135     struct Qdisc *child = q->qdisc;
136
137     return child->ops->peek(child);
138 }
139
140 static void red_reset(struct Qdisc *sch)
141 {
142     struct red_sched_data *q = qdisc_priv(sch);
143
144     qdisc_reset(q->qdisc);
145     sch->qstats.backlog = 0;
146     sch->q.qlen = 0;
147     red_restart(&q->vars);
148 }
149
150 static void red_destroy(struct Qdisc *sch)
151 {
152     struct red_sched_data *q = qdisc_priv(sch);
153
154     del_timer_sync(&q->adapt_timer);
155     qdisc_destroy(q->qdisc);
156 }
157
158 static const struct nla_policy red_policy[TCA_RED_MAX + 1] = {
159     [TCA_RED_PARMS] = { .len = sizeof(struct tc_red_qopt) },
160     [TCA_RED_STAB] = { .len = RED_STAB_SIZE },
161     [TCA_RED_MAX_P] = { .type = NLA_U32 },
162 };
163
164 static int red_change(struct Qdisc *sch, struct nlattr *opt)
165 {
166     struct red_sched_data *q = qdisc_priv(sch);
167     struct nlattr *tb[TCA_RED_MAX + 1];
168     struct tc_red_qopt *ctl;
169     struct Qdisc *child = NULL;
170     int err;
171     u32 max_p;

```

```

172
173     if (opt == NULL)
174         return -EINVAL;
175
176     err = nla_parse_nested(tb, TCA_RED_MAX, opt, red_policy);
177     if (err < 0)
178         return err;
179
180     if (tb[TCA_RED_PARAMS] == NULL ||
181         tb[TCA_RED_STAB] == NULL)
182         return -EINVAL;
183
184     max_P = tb[TCA_RED_MAX_P] ? nla_get_u32(tb[TCA_RED_MAX_P]) : 0;
185
186     ctl = nla_data(tb[TCA_RED_PARAMS]);
187
188     if (ctl->limit > 0) {
189         child = fifo_create_dflt(sch, &bififo_qdisc_ops, ctl->limit);
190         if (IS_ERR(child))
191             return PTR_ERR(child);
192     }
193
194     sch_tree_lock(sch);
195     q->flags = ctl->flags;
196     q->limit = ctl->limit;
197     if (child) {
198         qdisc_tree_reduce_backlog(q->qdisc, q->qdisc->q.qlen,
199                                 q->qdisc->qstats.backlog);
200         qdisc_destroy(q->qdisc);
201         q->qdisc = child;
202     }
203
204     red_set_parms(&q->parms,
205                  ctl->qth_min, ctl->qth_max, ctl->Wlog,
206                  ctl->Plog, ctl->Scell_log,
207                  nla_data(tb[TCA_RED_STAB]),
208                  max_P);
209     red_set_vars(&q->vars);
210
211     del_timer(&q->adapt_timer);
212     if (ctl->flags & TC_RED_ADAPTATIVE)
213         mod_timer(&q->adapt_timer, jiffies + HZ/2);
214
215     if (!q->qdisc->q.qlen)
216         red_start_of_idle_period(&q->vars);
217
218     sch_tree_unlock(sch);
219     return 0;
220 }
221
222 static inline void red_adaptative_timer(unsigned long arg)
223 {
224     struct Qdisc *sch = (struct Qdisc *)arg;
225     struct red_sched_data *q = qdisc_priv(sch);
226     spinlock_t *root_lock = qdisc_lock(qdisc_root_sleeping(sch));
227
228     spin_lock(root_lock);
229     red_adaptative_algo(&q->parms, &q->vars);
230     mod_timer(&q->adapt_timer, jiffies + HZ/2);
231     spin_unlock(root_lock);
232 }
233
234 static int red_init(struct Qdisc *sch, struct nlattr *opt)
235 {
236     struct red_sched_data *q = qdisc_priv(sch);
237
238     q->qdisc = &noop_qdisc;
239     setup_timer(&q->adapt_timer, red_adaptative_timer, (unsigned long)sch);
240     return red_change(sch, opt);
241 }
242
243 static int red_dump(struct Qdisc *sch, struct sk_buff *skb)
244 {
245     struct red_sched_data *q = qdisc_priv(sch);
246     struct nlattr *opts = NULL;
247     struct tc_red_qopt opt = {
248         .limit      = q->limit,
249         .flags      = q->flags,
250         .qth_min    = q->parms.qth_min >> q->parms.Wlog,
251         .qth_max    = q->parms.qth_max >> q->parms.Wlog,
252         .Wlog       = q->parms.Wlog,
253         .Plog       = q->parms.Plog,
254         .Scell_log  = q->parms.Scell_log,
255     };
256
257     sch->qstats.backlog = q->qdisc->qstats.backlog;
258     opts = nla_nest_start(skb, TCA_OPTIONS);
259     if (opts == NULL)
260         goto nla_put_failure;
261     if (nla_put(skb, TCA_RED_PARAMS, sizeof(opt), &opt) ||
262         nla_put_u32(skb, TCA_RED_MAX_P, q->parms.max_P))
263         goto nla_put_failure;
264     return nla_nest_end(skb, opts);
265
266 nla_put_failure:

```

```

267     nla_nest_cancel(skb, opts);
268     return -EMSGSIZE;
269 }
270
271 static int red_dump_stats(struct Qdisc *sch, struct gnet_dump *d)
272 {
273     struct red_sched_data *q = qdisc_priv(sch);
274     struct tc_red_xstats st = {
275         .early = q->stats.prob_drop + q->stats.forced_drop,
276         .pdrop = q->stats.pdrop,
277         .other = q->stats.other,
278         .marked = q->stats.prob_mark + q->stats.forced_mark,
279     };
280
281     return gnet_stats_copy_app(d, &st, sizeof(st));
282 }
283
284 static int red_dump_class(struct Qdisc *sch, unsigned long cl,
285                          struct sk_buff *skb, struct tcmsg *tcm)
286 {
287     struct red_sched_data *q = qdisc_priv(sch);
288
289     tcm->tcm_handle |= TC_H_MIN(1);
290     tcm->tcm_info = q->qdisc->handle;
291     return 0;
292 }
293
294 static int red_graft(struct Qdisc *sch, unsigned long arg, struct Qdisc *new,
295                    struct Qdisc **old)
296 {
297     struct red_sched_data *q = qdisc_priv(sch);
298
299     if (new == NULL)
300         new = &noop_qdisc;
301
302     *old = qdisc_replace(sch, new, &q->qdisc);
303     return 0;
304 }
305
306 static struct Qdisc *red_leaf(struct Qdisc *sch, unsigned long arg)
307 {
308     struct red_sched_data *q = qdisc_priv(sch);
309     return q->qdisc;
310 }
311
312 static unsigned long red_get(struct Qdisc *sch, u32 classid)
313 {
314     return 1;
315 }
316
317 static void red_put(struct Qdisc *sch, unsigned long arg)
318 {
319 }
320
321 static void red_walk(struct Qdisc *sch, struct qdisc_walker *walker)
322 {
323     if (!walker->stop) {
324         if (walker->count >= walker->skip)
325             if (walker->fn(sch, 1, walker) < 0) {
326                 walker->stop = 1;
327                 return;
328             }
329         walker->count++;
330     }
331 }
332
333 static const struct Qdisc_class_ops red_class_ops = {
334     .graft = red_graft,
335     .leaf = red_leaf,
336     .get = red_get,
337     .put = red_put,
338     .walk = red_walk,
339     .dump = red_dump_class,
340 };
341
342 static struct Qdisc_ops red_qdisc_ops __read_mostly = {
343     .id = "red",
344     .priv_size = sizeof(struct red_sched_data),
345     .cl_ops = &red_class_ops,
346     .enqueue = red_enqueue,
347     .dequeue = red_dequeue,
348     .peek = red_peek,
349     .init = red_init,
350     .reset = red_reset,
351     .destroy = red_destroy,
352     .change = red_change,
353     .dump = red_dump,
354     .dump_stats = red_dump_stats,
355     .owner = THIS_MODULE,
356 };
357
358 static int __init red_module_init(void)
359 {
360     return register_qdisc(&red_qdisc_ops);
361 }

```

```
362
363 static void __exit red_module_exit(void)
364 {
365     unregister_qdisc(&red_qdisc_ops);
366 }
367
368 module_init(red_module_init)
369 module_exit(red_module_exit)
370
371 MODULE_LICENSE("GPL");
372
```

This page was automatically generated by LXR 0.3.1 (source). • Linux is a registered trademark of Linus Torvalds • Contact us

[HOME](#) [DEVELOPMENT](#) [SERVICES](#) [TRAINING](#) [DOCS](#) [COMMUNITY](#) [COMPANY](#) [BLOG](#)